

---

# Data Cleaning

---

## Imports

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.under_sampling import RandomUnderSampler
from sklearn.preprocessing import StandardScaler
```

---

## Loading Original Data

```
In [ ]: df = pd.read_csv('weatherAUS.csv')
df.head()
```

```
Out[ ]:
```

|   | Date       | Location | MinTemp | MaxTemp | Rainfall | Evaporation | Sunshine | WindGustDir | Wind |
|---|------------|----------|---------|---------|----------|-------------|----------|-------------|------|
| 0 | 2008-12-01 | Albury   | 13.4    | 22.9    | 0.6      | NaN         | NaN      | W           |      |
| 1 | 2008-12-02 | Albury   | 7.4     | 25.1    | 0.0      | NaN         | NaN      | WNW         |      |
| 2 | 2008-12-03 | Albury   | 12.9    | 25.7    | 0.0      | NaN         | NaN      | WSW         |      |
| 3 | 2008-12-04 | Albury   | 9.2     | 28.0    | 0.0      | NaN         | NaN      | NE          |      |
| 4 | 2008-12-05 | Albury   | 17.5    | 32.3    | 1.0      | NaN         | NaN      | W           |      |

5 rows x 23 columns

---

## Binary Encoding of RainToday/Tomorrow

```
In [ ]: # One Hot Encoding

# Update RainToday column
raintoday = pd.get_dummies(df['RainToday'])
df['RainToday'] = raintoday['Yes']

# Update RainTomorrow column
raintomorrow = pd.get_dummies(df['RainTomorrow'])
df['RainTomorrow'] = raintomorrow['Yes']

df.head()
```

```
Out[ ]:
```

|   | Date       | Location | MinTemp | MaxTemp | Rainfall | Evaporation | Sunshine | WindGustDir | Wind |
|---|------------|----------|---------|---------|----------|-------------|----------|-------------|------|
| 0 | 2008-12-01 | Albury   | 13.4    | 22.9    | 0.6      | NaN         | NaN      | W           |      |
| 1 | 2008-12-02 | Albury   | 7.4     | 25.1    | 0.0      | NaN         | NaN      | WNW         |      |
| 2 | 2008-12-03 | Albury   | 12.9    | 25.7    | 0.0      | NaN         | NaN      | WSW         |      |
| 3 | 2008-12-04 | Albury   | 9.2     | 28.0    | 0.0      | NaN         | NaN      | NE          |      |
| 4 | 2008-12-05 | Albury   | 17.5    | 32.3    | 1.0      | NaN         | NaN      | W           |      |

5 rows × 23 columns

---

## Remove NaNs

```
In [ ]: # Remove columns with over 1/3 NaNs
df.drop(['Evaporation', 'Sunshine', 'Rainfall', 'Cloud9am', 'Cloud3pm', 'Loc

df.head()
```

| Out[ ]: |  | Date       | MinTemp | MaxTemp | WindGustDir | WindGustSpeed | WindDir9am | WindDir3pm | Wi |
|---------|--|------------|---------|---------|-------------|---------------|------------|------------|----|
| 0       |  | 2008-12-01 | 13.4    | 22.9    | W           | 44.0          | W          | WNW        |    |
| 1       |  | 2008-12-02 | 7.4     | 25.1    | WNW         | 44.0          | NNW        | WSW        |    |
| 2       |  | 2008-12-03 | 12.9    | 25.7    | WSW         | 46.0          | W          | WSW        |    |
| 3       |  | 2008-12-04 | 9.2     | 28.0    | NE          | 24.0          | SE         | E          |    |
| 4       |  | 2008-12-05 | 17.5    | 32.3    | W           | 41.0          | ENE        | NW         |    |

```
In [ ]: # Drop NaN values
df.dropna(inplace = True)

df.head()
```

| Out[ ]: |  | Date       | MinTemp | MaxTemp | WindGustDir | WindGustSpeed | WindDir9am | WindDir3pm | Wi |
|---------|--|------------|---------|---------|-------------|---------------|------------|------------|----|
| 0       |  | 2008-12-01 | 13.4    | 22.9    | W           | 44.0          | W          | WNW        |    |
| 1       |  | 2008-12-02 | 7.4     | 25.1    | WNW         | 44.0          | NNW        | WSW        |    |
| 2       |  | 2008-12-03 | 12.9    | 25.7    | WSW         | 46.0          | W          | WSW        |    |
| 3       |  | 2008-12-04 | 9.2     | 28.0    | NE          | 24.0          | SE         | E          |    |
| 4       |  | 2008-12-05 | 17.5    | 32.3    | W           | 41.0          | ENE        | NW         |    |

## Encode Remaining Categorical Variables

```
In [ ]: # Update Date Column

df['year'] = pd.DatetimeIndex(df['Date']).year
df['month'] = pd.DatetimeIndex(df['Date']).month
df['day'] = pd.DatetimeIndex(df['Date']).day
df.drop(['Date'], axis = 1, inplace = True)
df.head()
```

```
Out[ ]:
```

|   | MinTemp | MaxTemp | WindGustDir | WindGustSpeed | WindDir9am | WindDir3pm | WindSpeed |
|---|---------|---------|-------------|---------------|------------|------------|-----------|
| 0 | 13.4    | 22.9    | W           | 44.0          | W          | WNW        |           |
| 1 | 7.4     | 25.1    | WNW         | 44.0          | NNW        | WSW        |           |
| 2 | 12.9    | 25.7    | WSW         | 46.0          | W          | WSW        |           |
| 3 | 9.2     | 28.0    | NE          | 24.0          | SE         | E          |           |
| 4 | 17.5    | 32.3    | W           | 41.0          | ENE        | NW         |           |

```
In [ ]: df.WindGustDir.value_counts().sort_values(ascending=False)
```

```
Out[ ]:
```

|     |      |
|-----|------|
| W   | 8542 |
| SE  | 8215 |
| E   | 7995 |
| SSE | 7925 |
| S   | 7912 |
| N   | 7896 |
| WSW | 7821 |
| SW  | 7773 |
| SSW | 7522 |
| WNW | 6982 |
| ENE | 6904 |
| NW  | 6553 |
| ESE | 6315 |
| NE  | 6148 |
| NNE | 5604 |
| NNW | 5308 |

Name: WindGustDir, dtype: int64

```
In [ ]: def encoding_N(row):
    if 'N' in row[column]:
        return 1
    else:
        return 0

def encoding_E(row):
    if 'E' in row[column]:
        return 1
    else:
        return 0

def encoding_S(row):
    if 'S' in row[column]:
        return 1
    else:
        return 0

def encoding_W(row):
    if 'W' in row[column]:
        return 1
    else:
        return 0
```

```
In [ ]: update_columns = ['WindGustDir', 'WindDir9am', 'WindDir3pm']

for column in update_columns:
    df[column+'N'] = df.apply(lambda row: encoding_N(row), axis=1)
    df[column+'E'] = df.apply(lambda row: encoding_E(row), axis=1)
    df[column+'S'] = df.apply(lambda row: encoding_S(row), axis=1)
    df[column+'W'] = df.apply(lambda row: encoding_W(row), axis=1)

df.drop(update_columns, axis = 1, inplace = True)
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

|   | MinTemp | MaxTemp | WindGustSpeed | WindSpeed9am | WindSpeed3pm | Humidity9am | Hu |
|---|---------|---------|---------------|--------------|--------------|-------------|----|
| 0 | 13.4    | 22.9    | 44.0          | 20.0         | 24.0         | 71.0        |    |
| 1 | 7.4     | 25.1    | 44.0          | 4.0          | 22.0         | 44.0        |    |
| 2 | 12.9    | 25.7    | 46.0          | 19.0         | 26.0         | 38.0        |    |
| 3 | 9.2     | 28.0    | 24.0          | 11.0         | 9.0          | 45.0        |    |
| 4 | 17.5    | 32.3    | 41.0          | 7.0          | 20.0         | 82.0        |    |

5 rows × 28 columns

---

## Data Exploration

---

### Initial Descriptions

```
In [ ]: df.describe()
```

| Out [ ]:     | MinTemp       | MaxTemp       | WindGustSpeed | WindSpeed9am  | WindSpeed3pm  | Hu   |
|--------------|---------------|---------------|---------------|---------------|---------------|------|
| <b>count</b> | 115415.000000 | 115415.000000 | 115415.000000 | 115415.000000 | 115415.000000 | 1154 |
| <b>mean</b>  | 12.672407     | 23.644513     | 40.828081     | 15.216688     | 19.516371     |      |
| <b>std</b>   | 6.238970      | 6.971647      | 13.338630     | 8.362088      | 8.586030      |      |
| <b>min</b>   | -8.200000     | 2.600000      | 7.000000      | 2.000000      | 2.000000      |      |
| <b>25%</b>   | 8.100000      | 18.300000     | 31.000000     | 9.000000      | 13.000000     |      |
| <b>50%</b>   | 12.400000     | 23.100000     | 39.000000     | 13.000000     | 19.000000     |      |
| <b>75%</b>   | 17.200000     | 28.700000     | 48.000000     | 20.000000     | 24.000000     |      |
| <b>max</b>   | 33.900000     | 48.100000     | 135.000000    | 87.000000     | 87.000000     | 1    |


8 rows x 28 columns

## Correlation Matrix

```
In [ ]: # Correlation Matrix attributes with other attributes
```

```
corr_map = sns.heatmap(df.corr())
corr_map
```

```
Out [ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb552977910>
```



```
In [ ]: correlate = df.corr()
correlate
```

Out[ ]:

|                      | MinTemp   | MaxTemp   | WindGustSpeed | WindSpeed9am | WindSpeed3pm | Hi |
|----------------------|-----------|-----------|---------------|--------------|--------------|----|
| <b>MinTemp</b>       | 1.000000  | 0.728610  | 0.159282      | 0.150052     | 0.145654     |    |
| <b>MaxTemp</b>       | 0.728610  | 1.000000  | 0.056850      | -0.005989    | 0.013626     |    |
| <b>WindGustSpeed</b> | 0.159282  | 0.056850  | 1.000000      | 0.590132     | 0.678569     |    |
| <b>WindSpeed9am</b>  | 0.150052  | -0.005989 | 0.590132      | 1.000000     | 0.498255     |    |
| <b>WindSpeed3pm</b>  | 0.145654  | 0.013626  | 0.678569      | 0.498255     | 1.000000     |    |
| <b>Humidity9am</b>   | -0.217067 | -0.515588 | -0.183934     | -0.222205    | -0.093797    |    |
| <b>Humidity3pm</b>   | 0.026907  | -0.497169 | -0.021263     | -0.018744    | 0.052981     |    |
| <b>Pressure9am</b>   | -0.433600 | -0.312804 | -0.446087     | -0.202463    | -0.285172    |    |
| <b>Pressure3pm</b>   | -0.448056 | -0.411641 | -0.400743     | -0.151562    | -0.245085    |    |
| <b>Temp9am</b>       | 0.898502  | 0.884949  | 0.124632      | 0.089087     | 0.126241     |    |
| <b>Temp3pm</b>       | 0.703989  | 0.984276  | 0.018299      | -0.019171    | -0.011790    |    |
| <b>RainToday</b>     | 0.043948  | -0.240498 | 0.152262      | 0.097250     | 0.086494     |    |
| <b>RainTomorrow</b>  | 0.077364  | -0.167040 | 0.234628      | 0.091620     | 0.095698     |    |
| <b>year</b>          | 0.045828  | 0.062946  | -0.027563     | -0.014091    | -0.027539    |    |
| <b>month</b>         | -0.211504 | -0.171276 | 0.051801      | 0.042637     | 0.057476     |    |
| <b>day</b>           | 0.001575  | -0.001384 | -0.008812     | -0.006671    | -0.009200    |    |
| <b>WindGustDirN</b>  | -0.006184 | 0.058115  | 0.003732      | -0.039966    | 0.004605     |    |
| <b>WindGustDirE</b>  | 0.236121  | 0.239784  | -0.167436     | -0.022673    | -0.092648    |    |
| <b>WindGustDirS</b>  | 0.011241  | -0.064723 | -0.021431     | 0.028567     | 0.001690     |    |
| <b>WindGustDirW</b>  | -0.208689 | -0.203830 | 0.140730      | -0.021348    | 0.071292     |    |
| <b>WindDir9amN</b>   | -0.082341 | 0.017123  | 0.094629      | -0.021269    | 0.034362     |    |
| <b>WindDir9amE</b>   | 0.150988  | 0.294265  | -0.125451     | -0.033222    | -0.133435    |    |
| <b>WindDir9amS</b>   | 0.090011  | -0.025456 | -0.086797     | 0.010572     | -0.011101    |    |
| <b>WindDir9amW</b>   | -0.145901 | -0.283775 | 0.062170      | -0.020491    | 0.069909     |    |
| <b>WindDir3pmN</b>   | 0.020486  | 0.101639  | 0.024387      | -0.064041    | -0.001086    |    |
| <b>WindDir3pmE</b>   | 0.232103  | 0.175630  | -0.157909     | -0.023885    | -0.074840    |    |
| <b>WindDir3pmS</b>   | -0.019005 | -0.102426 | -0.033484     | 0.063355     | -0.004145    |    |
| <b>WindDir3pmW</b>   | -0.197892 | -0.138989 | 0.144229      | -0.014892    | 0.064907     |    |

28 rows x 28 columns

```
In [ ]: correlated_features = set()

for i in range(len(correlate.columns)):
    for j in range(i):
        if abs(correlate.iloc[i, j]) > 0.85:
            colname = correlate.columns[i]
            correlated_features.add(colname)

df.drop(labels=correlated_features, axis=1, inplace=True)
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

|   | MinTemp | MaxTemp | WindGustSpeed | WindSpeed9am | WindSpeed3pm | Humidity9am | Hu |
|---|---------|---------|---------------|--------------|--------------|-------------|----|
| 0 | 13.4    | 22.9    | 44.0          | 20.0         | 24.0         | 71.0        |    |
| 1 | 7.4     | 25.1    | 44.0          | 4.0          | 22.0         | 44.0        |    |
| 2 | 12.9    | 25.7    | 46.0          | 19.0         | 26.0         | 38.0        |    |
| 3 | 9.2     | 28.0    | 24.0          | 11.0         | 9.0          | 45.0        |    |
| 4 | 17.5    | 32.3    | 41.0          | 7.0          | 20.0         | 82.0        |    |

5 rows × 25 columns

---

## Balancing Our Data

---

### Undersampling

```
In [ ]: # Barplot to display original balance of data

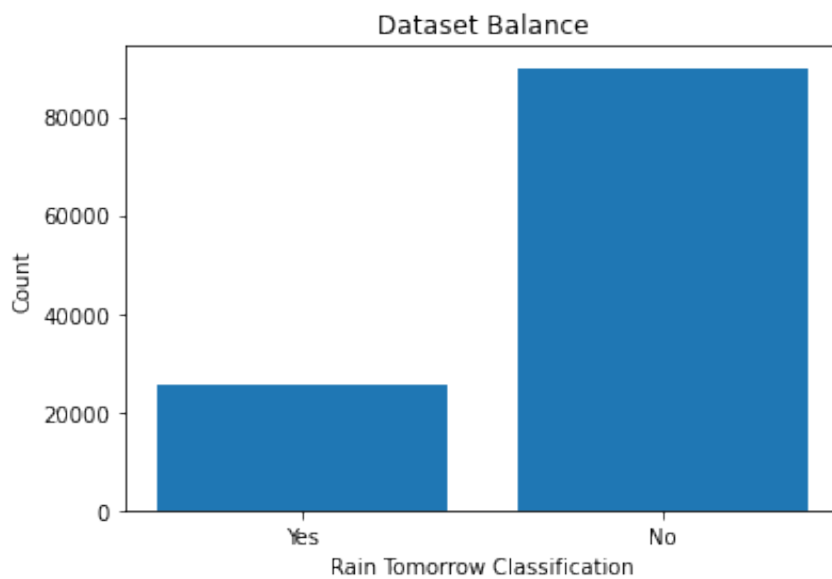
sum_yes = df['RainTomorrow'].sum()
sum_no = len(df) - sum_yes

x_labels = ['Yes', 'No']
y_labels = [sum_yes, sum_no]

plt.bar(x_labels, y_labels)
plt.xlabel('Rain Tomorrow Classification')
plt.ylabel('Count')
plt.title('Dataset Balance')
```

```
Out[ ]: Text(0.5, 1.0, 'Dataset Balance')
```





```
In [ ]: # Undersampling
X = df.drop("RainTomorrow", axis=1)
y = df["RainTomorrow"]
RUS = RandomUnderSampler(random_state=42)
X_und, y_und = RUS.fit_sample(X, y)
```

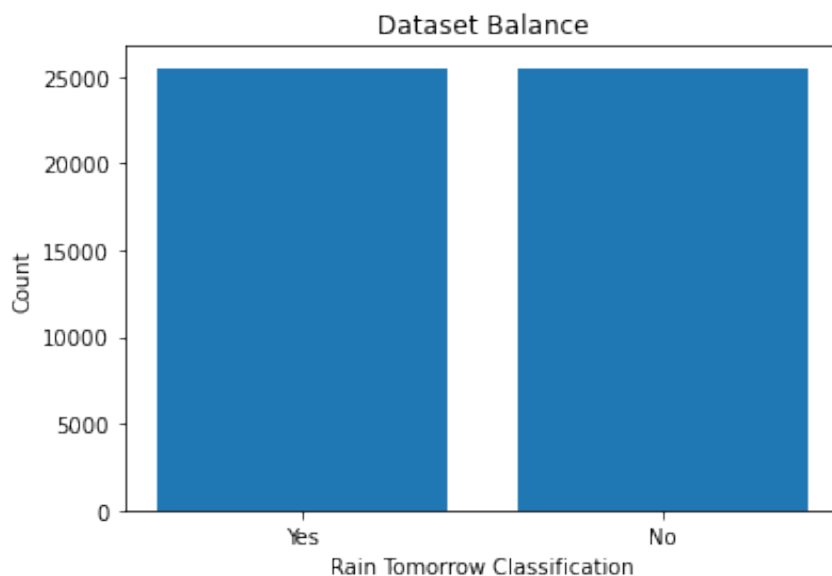
```
In [ ]: # Barplot of data after undersampling

sum_yes = y_und.sum()
sum_no = len(y_und) - sum_yes

x_labels = ['Yes', 'No']
y_labels = [sum_yes, sum_no]

plt.bar(x_labels, y_labels)
plt.xlabel('Rain Tomorrow Classification')
plt.ylabel('Count')
plt.title('Dataset Balance')
```

```
Out[ ]: Text(0.5, 1.0, 'Dataset Balance')
```



---

## Standardize Data

---

```
In [ ]: scaler = StandardScaler()  
        scaler.fit(X_und)  
        standardized_X = scaler.transform(features)
```

---

## Download the New Dataset

---

```
In [ ]: # Download  
        standardized_X.to_csv("cleaned_rain_x.csv")  
        y_und.to_csv("cleaned_rain_y.csv")
```